

ACID PROPERTIES

By Pavani Korada

ACID


ACID stands for:

- Atomicity
- Consistency
- Isolation
- Durability


These are the four essential properties that make sure database transactions are correct, complete, and reliable.



ATOMICITY




- Ensures a transaction is treated as a single unit.
- Either all operations succeed, or none do.
- No partial changes are left if something fails midway.
- **Example:** transferring money from one account to another must debit and credit together, or not happen at all.





CONSISTENCY

Consistency


Preserving database invariants



- After a transaction, the database must stay in a valid state.
- All rules (constraints, data types, relationships) must still hold.
- If a transaction breaks a rule, it gets rolled back.
- This keeps data “correct” before and after transactions.




ISOLATION



- Even when many transactions run at the same time, each one behaves as if it runs alone.
- One transaction should not see half-done changes from another.
- This prevents issues like dirty reads and inconsistent results.


Example:

- Two users updating the same record at the same time should not corrupt the data.
- Databases use locking or version control to handle this.




ISOLATION LEVELS



Isolation Level	Violations		
	Dirty Read	Non-repeatable Read	Phantom Read
Serializable	Don't occur	Don't occur	Don't occur
Repeatable	Don't occur	Don't occur	May occur
Read committed	Don't occur	May occur	May occur



A dirty read happens when a transaction reads data modified by another transaction that hasn't been committed yet.



This happens when a transaction reads the same data twice and gets different results because another transaction modified the data in between.




A phantom read occurs when a transaction re-runs a query and gets additional or missing rows because another transaction inserted or deleted matching rows.

Repeatable Read prevents non-repeatable reads. It gives each transaction a consistent snapshot of the data, but Phantom reads can still occur.

DURABILITY

Durability


Data is persisted after transaction is committed even in a system failure



Once a transaction is committed, it stays committed — even if:

- System crashes
- Power goes off
- Server restarts


The data is permanently stored. This guarantees that completed work isn't lost.
This ensures reliability.



- Writing transaction logs
- Using Write-Ahead Logging (WAL)
- Persisting changes to disk before confirming commit

In distributed databases:

- Data is replicated across multiple nodes
- If one node fails, committed transactions are safe on other nodes



WHY ACID IS IMPORTANT

As a data engineer, we deal with:

- ETL/ELT pipelines
- Data ingestion
- Batch & streaming jobs
- Data warehousing



Without
ACID



- Pipelines may create duplicate data
- Partial writes may happen
- Corrupted tables can break dashboards
- Analytics results become unreliable

HOW DATABRICKS USES ACID

- Atomic Write
 - If a Spark job fails while writing data to a Delta table, the entire write is rolled back.
- Consistency with Schema Enforcement
 - Delta Lake enforces schema rules. If you try to insert wrong columns, it throws an error.
- Isolation with Optimistic Concurrency Control
 - Multiple users can update the same Delta table safely.
 - Databricks handles conflicts automatically.
- Durability with Transaction Log
 - Delta Lake stores a transaction log (_delta_log folder).
 - This ensures:
 - Every change is tracked
 - Data can be restored
 - No data loss after commit



Thank
You